

# Codex Task Prompt (full, with appendices)

## Title

Execution hardening: stabilize pending-order handling (no oscillation), reduce duplicate intents, improve pacing-cap semantics, enrich pending-order telemetry (counts/age/samples + transition-based warnings), fix noisy model-liveness alerts (phase-aware + diagnosable), and remove redundant env knobs in code (canonical keys only, fail-fast on deprecated)

## Repository Context

Repo: ai-trading-bot (Dom)

Primary runtime: systemd service ai-trading.service on Ubuntu droplet using .env and generated .env.runtime.

Key modules involved:

- ai\_trading/core/bot\_engine.py (cycle orchestration, pending order gating, symbol blocking, signal emission hooks)
- ai\_trading/execution/live\_trading.py (execution engine, pending-new timeout policy, pacing cap, duplicate intent gate)
- ai\_trading/main.py (cycle loop, alert emission including ALERT\_MODEL\_LIVENESS)
- ai\_trading/monitoring/model\_liveness.py (liveness evaluation, snapshot helpers)
- ai\_trading/position\_sizing.py + ai\_trading/config/runtime.py + ai\_trading/settings.py (env resolution + deprecated keys)
- ai\_trading/utils/env.py, ai\_trading/diagnostics/env\_diag.py, ai\_trading/broker/alpaca\_credentials.py, ai\_trading/validation/validate\_env.py (Alpaca env usage + aliases)

Tests (existing targets likely touched):

- tests/bot\_engine/\*pending\*

- `tests/execution/test_execution_runtime_controls.py`
- `tests/execution/test_order_pacing_cap.py`
- plus new tests described below

Observed operational symptoms (`journalctl`):

- Frequent `PENDING_ORDERS_STILL_PRESENT` warnings and repeated `PENDING_ORDERS_POLICY_APPLIED`
- Block decay “release” can occur while orders are still open (`PENDING_SYMBOL_BLOCK_DECAY_RELEASED` then `PENDING_ORDERS_STILL_PRESENT`)
- Regular `cycle_duplicate_intent`
- `ORDER_SUBMIT_SKIPPED_DETAIL` with `reason=order_pacing_cap` happens often
- `ALERT_MODEL_LIVENESS` fires constantly as `ERROR`, creating noise that can hide real faults

Operator debugging needs (must address):

When `PENDING_ORDERS_STILL_PRESENT` fires, logs must include:

- `open_count`
- `pending_count` (or whatever internal classification is used)
- `oldest_open_age_s`
- bounded sample\_orders (ids/statuses + timestamps/age)

And ideally:

- `WARN` only on backlog state transitions (start), then periodic heartbeat every N minutes while backlog persists.

Droplet findings to fold in (new):

- `.env` and `service env` have been cleaned so legacy keys like `ALPACA_BASE_URL`, `AI_TRADING_MAX_POSITION_SIZE`, `AI_TRADING_MAX_DRAWDOWN_THRESHOLD`, and `NEWS_API_KEY` are no longer present. Code must now **stop accepting deprecated keys silently** and either (a) fail-fast with actionable message or (b) keep only canonical reads (NO SHIMS).
- Multiple Alpaca URL keys still exist and are being referenced across code; the droplet currently sets:

- ALPACA\_TRADING\_BASE\_URL=https://paper-api.alpaca.markets
- ALPACA\_API\_URL=https://paper-api.alpaca.markets
- ALPACA\_BASE\_URL=https://paper-api.alpaca.markets (legacy)
- ALPACA\_DATA\_BASE\_URL=https://data.alpaca.markets
- CLI operator attempt python -m ai\_trading --print-config failed (“unrecognized arguments”), so there is currently no first-class “print resolved config” pathway.

## Problem Statement

Pending-order handling is safe, but:

- Too noisy (warnings for normal open-limit behavior)
- Oscillatory (block decay can release while orders persist → reblock next cycle)
- Inefficient (duplicate intents leak into execution)
- Over-coupled to pacing (maintenance/duplicates contribute to “new order” pacing cap)

Additionally, model liveness alerting is broken operationally:

- ALERT\_MODEL\_LIVENESS spams ERROR repeatedly (appears ~every few minutes), likely because liveness is evaluated when signals are not expected, or signal “note” functions are not called, or liveness ignores execution phases / warmup / market-closed logic.
- Alerts are not diagnosable (no structured payload, and snapshot isn’t surfaced in health/logs).

Finally, env redundancy is still present in code paths:

- Deprecated keys are still read in leaf modules (e.g., ai\_trading/position\_sizing.py reading deprecated AI\_TRADING\_MAX\_POSITION\_SIZE).
- Alpaca base URL keys exist in multiple forms; this makes .env harder to reason about and creates inconsistency risk.

## Scope of Work

### A) Make pending-new timeout policy idempotent per engine cycle

**Goal:** Avoid running pending-new timeout policy multiple times in one engine cycle.

In `ai_trading/execution/live_trading.py`:

- Add guard so `_apply_pending_new_timeout_policy()` runs at most once per `engine_cycle_index`.
- Track `self._pending_new_policy_last_cycle_index: int | None`.
- If equal to current cycle index, return early.
- Set it when policy runs.
- Ensure tests can call `_apply_pending_new_timeout_policy()` directly:
  - Provide a safe way to reset internal state in tests (e.g., a private reset helper in the class and only used by tests).

### B) Fix block decay oscillation: never “decay release” while orders are still open

**Goal:** Stop `release→reblock` loop.

In `ai_trading/core/bot_engine.py`:

- Only release a symbol block due to decay if there are **zero open orders** for that symbol.
- If open orders still exist, keep the block and log:
  - `PENDING_SYMBOL_BLOCK_DECAY_DEFERRED` with structured payload:
    - `symbol`
    - `open_orders_count`
    - `oldest_open_order_age_s`

## C) Reduce noise + add telemetry for pending/open orders (structured + severity gated)

**Goal:** Make warnings meaningful, and make backlog diagnosable.

In ai\_trading/core/bot\_engine.py pending-order gating/logging path:

### *C1) Compute age fields robustly*

- Compute oldest\_open\_age\_s using broker timestamps:
  - prefer updated\_at, fallback submitted\_at
- Parse defensively; never crash if timestamps missing/unparseable (set fields to None).

### *C2) Add warning threshold env var*

- Add env var: AI\_TRADING\_PENDING\_ORDERS\_WARN\_AFTER\_SEC (default 120)
- If open orders exist but oldest\_open\_age\_s < warn\_after\_sec: log at INFO (or DEBUG), not WARNING.
- WARNING only when older than threshold and policy action is being applied (or backlog “stale”).

### *C3) Enrich pending-order payload (required)*

Update PENDING\_ORDERS\_DETECTED / PENDING\_ORDERS\_STILL\_PRESENT payload to include:

- open\_count (explicit definition: broker-active orders, not filled/canceled/expired)
- pending\_count (explicit definition: your “not yet stable/ack” set; document which statuses count)
- counts\_by\_status (status → count)
- oldest\_open\_age\_s
- affected\_symbols\_count
- sample\_orders (bounded, N=3 or 5 max), each:
  - order\_id (may be abbreviated but must be correlatable)
  - symbol
  - status

- submitted\_at and/or updated\_at (optional)
- age\_s (optional)
- Sampling rule: pick **oldest orders first**.

#### *C4) Transition-based warnings + periodic heartbeat (preferred)*

Implement stateful backlog emission:

- Maintain:
  - self.\_pending\_backlog\_active: bool
  - self.\_pending\_backlog\_last\_warn\_ts: float | None
- On transition False -> True:
  - Emit PENDING\_ORDERS\_BACKLOG\_STARTED (or reuse PENDING\_ORDERS\_STILL\_PRESENT with transition="started")
  - Severity:
    - WARNING only if oldest\_open\_age\_s >= warn\_after\_sec
    - else INFO
- While backlog persists:
  - Emit WARNING at most once per AI\_TRADING\_PENDING\_ORDERS\_WARN\_EVERY\_SEC (default 300)
  - Otherwise remain silent (or optional INFO heartbeat with cooldown)
- On transition True -> False:
  - Emit PENDING\_ORDERS\_CLEARED with summary payload.

Add env var:

- AI\_TRADING\_PENDING\_ORDERS\_WARN\_EVERY\_SEC (default 300)

Hard requirement: do not spam per-cycle.

#### **D) Dedupe cycle intents upstream (reduce cycle\_duplicate\_intent)**

**Goal:** Prevent duplicates from ever reaching execution.

In ai\_trading/core/bot\_engine.py (or wherever the execution candidate list is built):

- Dedupe intents by (symbol, side) before calling execution engine.
- Keep “best” candidate (deterministic tiebreaker):
  - primary: highest score (or strongest signal metric)
  - secondary: highest notional (or stable metric)
  - tertiary: stable ordering key (symbol, side, candidate\_id)
- Emit one bounded summary log per cycle:
  - CYCLE\_INTENTSDEDUPED payload: kept\_count, dropped\_count, dropped\_examples (cap small N)

Execution engine keeps existing cycle\_duplicate\_intent guard as last line of defense.

## E) Improve pacing-cap semantics (maintenance should not consume “new order” budget)

**Goal:** Pacing cap measures new order submissions, not cancel/replace maintenance.

In ai\_trading/execution/live\_trading.py:

- Split counters:
  - new\_orders\_submitted\_this\_cycle
  - maintenance\_actions\_this\_cycle (cancel/replace)
- Ensure pending-new policy maintenance actions do not increment EXECUTION\_MAX\_NEW\_ORDERS\_PER\_CYCLE counter.
- Update ORDER\_PACING\_CAP\_HIT / ORDER\_SUBMIT\_SKIPPED\_DETAIL payload so it clearly indicates:
  - cap\_type=new\_orders vs cap\_type=maintenance (if maintenance gets separate cap)
  - include limit, used, and headroom if available

Tests must confirm the split.

## NEW: F) Remove redundant env knobs in code (canonical keys only; NO SHIMS; fail-fast on deprecated keys)

**Principle:** pick one canonical env var per concept. Code reads only canonical key. If deprecated key is present, **fail fast** with actionable rename instructions (no silent aliases).

### F1) Remove MAX\_POSITION\_SIZE / AI\_TRADING\_MAX\_POSITION\_SIZE split for real

Observed:

- Config/runtime indicates AI\_TRADING\_MAX\_POSITION\_SIZE is deprecated, but ai\_trading/position\_sizing.py still reads it.

PR change:

- Update ai\_trading/position\_sizing.py to read **only** MAX\_POSITION\_SIZE.
- If AI\_TRADING\_MAX\_POSITION\_SIZE is set:
  - do **not** silently accept it
  - raise a startup error or emit a single high-visibility WARNING and exit (depending on your existing “strict readiness”/startup validation behavior), telling operator to rename to MAX\_POSITION\_SIZE.

Update tests accordingly.

### F2) Deduplicate other redundant knobs (repo scan + normalize)

Perform a repo-wide scan for env pairs where both old/new keys are being read (examples already found):

- MAX\_DRAWDOWN\_THRESHOLD vs AI\_TRADING\_MAX\_DRAWDOWN\_THRESHOLD
- AI\_TRADING\_ALLOW\_SHORT vs TRADING\_\_ALLOW\_SHORTS
- EXECUTION\_ALLOW\_FALLBACK\_WITHOUT\_NBBO vs AI\_TRADING\_EXEC\_ALLOW\_FALLBACK\_WITHOUT\_NBBO

- SENTIMENT\_API\_KEY vs NEWS\_API\_KEY (your .env currently sets both to the same value; code has an alias relationship)
- Any other deprecated\_env={...} entries in ai\_trading/config/runtime.py where leaf modules still read the deprecated key directly

PR requirement:

- Choose canonical key for each pair (prefer what ai\_trading/config/runtime.py declares canonical).
- Update leaf modules to read only canonical key.
- If deprecated keys are present: fail-fast with actionable message (NO SHIMS).
- Add a short “Env Canonicalization Map” doc/comment listing canonical key → deprecated keys and the error message behavior.

## NEW: F3) Canonicalize Alpaca URL environment variables in code (trading vs data)

Right now the repo references Alpaca endpoints across multiple keys, and ops ends up setting 2-3 keys to the same value to stay safe.

**Goal:** Make trading base URL and data base URL explicit, reduce redundancy, and stop reading legacy keys.

### Canonical keys

- Canonical trading/orders/account host: ALPACA\_TRADING\_BASE\_URL
- Canonical data host: ALPACA\_DATA\_BASE\_URL

### Deprecated keys to eliminate from *code reads*

- ALPACA\_API\_URL (currently widely used)
- ALPACA\_BASE\_URL (legacy)

PR behavior:

- All trading client construction must use ALPACA\_TRADING\_BASE\_URL.

- All data client construction must use ALPACA\_DATA\_BASE\_URL.
- If code currently expects ALPACA\_API\_URL / ALPACA\_BASE\_URL, migrate it.
- If deprecated keys are set at runtime:
  - do **not** silently accept
  - fail fast with a clear message:
    - “Set ALPACA\_TRADING\_BASE\_URL for trading endpoints; remove ALPACA\_API\_URL/ALPACA\_BASE\_URL”
- Update validation and diagnostics modules  
(ai\_trading/validation/validate\_env.py,  
ai\_trading/diagnostics/env\_diag.py, config management helpers) to reflect the new canonical rules.

**Important:** keep paper/live correctness checks:

- paper must be <https://paper-api.alpaca.markets>
- live must be <https://api.alpaca.markets> (or your configured live host)

(Enforce based on EXECUTION\_MODE or equivalent existing switch.)

## NEW: F4) Fix the operator “print config” / introspection gap

Observed on droplet:

- python -m ai\_trading --print-config fails (arg not recognized). Operators need a stable way to see resolved env/config, especially after canonicalization.

PR requirement:

- Add a supported CLI option (pick one):
  - python -m ai\_trading --print-config OR
  - python -m ai\_trading --diagnostics env / --dump-config
- It must:
  - Print resolved configuration **without secrets**
  - Show canonical key names (not deprecated ones)

- If deprecated keys are present, print the fail-fast message (and exit non-zero)

Add tests for this option if you have CLI parsing tests; otherwise add a small new test module that runs the argument parser.

## NEW: G) Fix ALERT\_MODEL\_LIVENESS so it only fires when meaningful (and becomes diagnosable)

Treat as a real bug.

### G1) Make liveness evaluation phase-aware

In ai\_trading/main.py / ai\_trading/monitoring/model\_liveness.py implement:

Liveness should be evaluated only when ALL are true:

1. Market is open (already passed in; must be respected).
2. Bot is in a phase where signals are expected:
  - a. post-bootstrap
  - b. post-reconcile / execution gate open
  - c. warmup complete
  - d. not in “market closed skip cycle”
3. Signals are expected given enabled modules:
  - a. If RL disabled (USE\_RL\_AGENT=0), don’t enforce RL heartbeat.
  - b. If ML disabled/model not loaded, don’t enforce ML heartbeat.
  - c. If execution is blocked for legitimate reasons, don’t alert liveness.

Implementation approach:

- In main.py, derive a boolean like signals\_expected\_now and pass into liveness evaluation (or compute inside liveness using shared state).
- In model\_liveness.py, if signals\_expected\_now is false: return no breaches.

## G2) Ensure signals are actually “noted” correctly

Likely root causes:

- `note_ml_signal` / `note_rl_signals_emitted` not called on the “real” signal emission path.
- Liveness expects a signal too frequently even when strategy legitimately produces none.

PR requirement:

- Identify where ML and RL signals are considered “emitted” in the actual flow.
- Decide and document semantics:
  - If “signal liveness” means “model produced a decision record (even no-trade)”, then note on decision production.
  - If it means “we produced at least one actionable candidate”, note only then.
- Enforce thresholds using:
  - `AI_TRADING_ML_SIGNAL_MAX_AGE_SECONDS`
  - `AI_TRADING_RL_SIGNAL_MAX_AGE_SECONDS`

## G3) Alert behavior: rate-limited + correct severity + structured payload

Acceptance:

- Rate-limit via `AI_TRADING_MODEL_LIVENESS_ALERT_COOLDOWN_SECONDS`
- Log level:
  - `WARNING` for breach
  - `ERROR` only if automated action taken (kill switch / canary rollback)

Every alert must include structured payload:

- `last_ml_signal_ts`, `last_rl_signal_ts` (or `None`)
- `ml_age_s`, `rl_age_s`
- `ml_max_age_s`, `rl_max_age_s`
- `market_open`
- `signals_expected_now`

- phase / execution gate state / warmup state (whatever exists)

#### G4) Add liveness snapshot to logs + health endpoint

`get_model_liveness_snapshot()` exists – wire it into:

- periodic health tick logs (cooldown; not per cycle)
- health endpoint payload (or diagnostics endpoint)

Snapshot must not leak secrets.

## Acceptance Criteria

Pending-orders / execution:

1. Pending-new timeout policy runs once per engine cycle (idempotent).
2. Symbol block decay does not release while open orders exist; emits `PENDING_SYMBOL_BLOCK_DECAY_DEFERRED`.
3. Pending backlog logs are severity-gated by age threshold:
  - a. `WARNING` only when `oldest_open_age_s >= AI_TRADING_PENDING_ORDERS_WARN_AFTER_SEC`
  - b. and warnings are transition/periodic, not per-cycle.
4. `PENDING_ORDERS_*` logs include:
  - a. `open_count`, `pending_count`, `counts_by_status`, `oldest_open_age_s`, bounded sample orders.
5. Upstream dedupe reduces `cycle_duplicate_intent` materially.
6. Pacing semantics split works:
  - a. maintenance actions do not consume new-order cap.

Env canonicalization:

7. `ai_trading/position_sizing.py` reads only `MAX_POSITION_SIZE`.
8. Deprecated env keys are not silently accepted:
  - a. if set, bot fails fast with clear rename instructions (NO SHIMS).
9. Similar dedupe performed for other redundant env pairs discovered.

10. Alpaca URLs:

- trading uses ALPACA\_TRADING\_BASE\_URL
- data uses ALPACA\_DATA\_BASE\_URL
- deprecated ALPACA\_API\_URL / ALPACA\_BASE\_URL are not read (fail-fast if present).

Operator introspection:

11. There is a supported CLI option to print resolved config (sanitized), including canonical env key names, and it exits non-zero if deprecated keys are set.

Model liveness:

12. ALERT\_MODEL\_LIVENESS no longer spams:

- only evaluated when signals expected and market open
- rate-limited
- WARNING unless action taken
- structured payload included

13. Health/logging includes a liveness snapshot so ops can debug quickly.

All tests pass.

## Change Details (Files / Touchpoints)

- ai\_trading/execution/live\_trading.py
  - idempotency guard for \_apply\_pending\_new\_timeout\_policy
  - split new-order pacing vs maintenance counters
  - improve ORDER\_PACING\_CAP\_HIT / ORDER\_SUBMIT\_SKIPPED\_DETAIL payload
- ai\_trading/core/bot\_engine.py
  - pending-order severity gating by age
  - telemetry enrichment (counts/age/samples)
  - state transition + heartbeat emission
  - block decay deferral
  - upstream intent dedupe
- ai\_trading/monitoring/model\_liveness.py

- phase-aware / signals-expected gating
  - cooldown correctness
  - structured payload + improved snapshot
- ai\_trading/main.py
  - derive signals\_expected\_now and apply gating
  - correct severity semantics (WARN vs ERROR on action)
  - periodic snapshot emission
- ai\_trading/position\_sizing.py
  - canonical MAX\_POSITION\_SIZE only; fail-fast if deprecated present
- Alpaca env canonicalization touchpoints (likely):
  - ai\_trading/utils/env.py
  - ai\_trading/broker/alpaca\_credentials.py
  - ai\_trading/diagnostics/env\_diag.py
  - ai\_trading/validation/validate\_env.py
  - any module that constructs Alpaca clients or uses base URLs
- CLI / diagnostics:
  - ai\_trading/\_\_main\_\_.py (argparse)
  - or a dedicated diagnostics module

## Tests (update/add)

Update or add tests to cover:

Pending orders:

- idempotency guard
- decay deferral
- telemetry payload fields present
- transition/heartbeat gating

Pacing cap:

- maintenance does not consume new-order cap

Env canonicalization:

- MAX\_POSITION\_SIZE works

- deprecated key present  $\Rightarrow$  fails fast with actionable message
- Alpaca deprecated URL key present  $\Rightarrow$  fails fast

Model liveness:

- does not evaluate/alert when market closed or phase not expecting signals
- alerts when market open + signals expected + last signal older than threshold
- includes payload fields
- respects cooldown
- WARNING vs ERROR semantics

CLI config dump:

- --print-config (or chosen flag) works, is sanitized, and honors fail-fast behavior

## Constraints & Standards

- **NO SHIMS** anywhere (no compatibility wrappers, no silent alias fallbacks).
- Fail-safe behavior (prefer blocking trading over unsafe repeated submissions).
- Logs structured and bounded; avoid per-cycle noise.
- Deterministic dedupe behavior.
- Use broker timestamps defensively; never crash on missing data.
- Preserve .env  $\rightarrow$  .env.runtime workflow.
- Do not delete .jsonl runtime logs (operator explicitly wants them retained).

## Validation Steps

```
cd /home/aiuser/ai-trading-bot  
source venv/bin/activate
```

```
pytest -q tests/bot_engine/test_order_pending_severity.py
pytest -q tests/bot_engine/test_pending_orders_cleanup.py
pytest -q tests/execution/test_execution_runtime_controls.py
pytest -q tests/execution/test_order_pacing_cap.py

# new tests you add (names TBD)
pytest -q tests/monitoring/test_model_liveness.py
pytest -q tests/config/test_env_canonicalization.py
pytest -q tests/cli/test_print_config.py

pytest -q
```

Operational check (droplet):

```
journalctl -u ai-trading.service --since "60 min ago" -o short-iso -l
| \
grep -E
'PENDING_ORDERS_|PENDING_SYMBOL_BLOCK_DECAY_|CYCLE_INTENTS_DEDUPED|cyc
le_duplicate_intent|order_pacing_cap|ORDER_SUBMIT_SKIPPED_DETAIL|ALERT
_MODEL_LIVENESS|MODEL_LIVENESS' | \
tail -n 400
```

Confirm:

- Pending backlog warnings are not spammy and include enriched payload.
- Liveness alerts are not constant; when they fire they're meaningful and include snapshot/payload.
- Deprecated env keys cause a clear, actionable startup failure (no silent fallback).
- Alpaca URL keys are canonicalized (trading vs data), and old keys are rejected.

## Risk & Rollback

Risk: pending-order semantics changes can alter how aggressively trading is blocked during open-limit workflows; liveness gating could suppress alerts if implemented incorrectly; env canonicalization could fail startup if operator forgets to rename a key.

Rollback:

- revert PR
- sudo systemctl restart ai-trading.service
- restore previous .env if needed (PR must make rename requirements obvious and explicit)

## Non-Goals

- No strategy logic changes (signals/scoring/model weights)
- No broker/provider rewrites
- No new compatibility shims or facades
- No deleting .jsonl runtime logs

## Appendices

### Appendix A: Target build/runtime combo (verbatim)

Ubuntu 24.04 (Noble) x86\_64, glibc 2.39, CPython 3.12.3 in venv at /home/aiuser/ai-trading-bot/venv. Pip compatible top tag: cp312-cp312-manylinux\_2\_39\_x86\_64 (and accepts older manylinux baselines). Dry-run shows runtime/dev reqs satisfied; notable heavy deps present: torch 2.3.1, stable-baselines3 2.3.2, gymnasium 0.29.1, scikit-learn 1.5.2, scipy 1.13.1, matplotlib 3.9.2, hypothesis 6.138.2. Alpaca SDK: alpaca-trade-api 3.2.0 (runtime) and alpaca-py 0.42.0 (dev-only).

## Appendix B: Current “known good” runtime knobs you’re using

- AI\_TRADING\_EXECUTION\_PHASE\_GATE\_ENABLED=1
- AI\_TRADING\_BLOCK\_ENTRIES\_ON\_FALLBACK\_MINUTE\_DATA=1
- AI\_TRADING\_EVENT\_DRIVEN\_NEW\_BAR\_ONLY=1
- EXECUTION\_MAX\_NEW\_ORDERS\_PER\_CYCLE=10
- Debug noise removed (AI\_TRADING\_EXEC\_DEBUG, AI\_TRADING\_EXEC\_LOG\_REQUESTS, etc.)

## Appendix C: Why this PR is worth it

- Less churn from block decay oscillation
- Cleaner logs (warnings mean “abnormal”, not “open orders exist”)
- Higher throughput under normal limit-order behavior
- Same safety posture, fewer self-inflicted throttles
- Better ops visibility: pending backlog payload includes counts/ages/samples
- Liveness monitoring becomes trustworthy: fewer false alarms, more diagnosable alerts, less masking of real issues
- Env becomes maintainable: canonical keys only, and the bot tells you exactly what to fix if a deprecated key appears